

Amazon Aurora MySQL Hands-on Lab – Version 2.2

Lab Guide

PART 1. Setup Environment.....	2
Task 1.1 – Sign in to Management Console, Select Region and Create Key Pair.....	2
Task 1.2 - Creating a Stack using CloudFormation	4
Task 1.3 - Connecting to the workstation EC2 instance	8
Task 1.4 – Set up the AWS CLI and seed the DB Cluster	8
PART 2. Cluster Endpoints and Auto Scaling.....	9
Task 2.1 – Running a read-only workload	9
PART 3. Cloning and backtracking databases	11
Task 3.1 - Creating a Clone	11
Task 3.2 - Backtracking the Database	14
PART 4. Performance Insights	15
Task 4.1 – Generate Load on Your Database Cluster	15
Part 5. Delete AWS resources	16
Task 5.1 – Delete Cloned Aurora Cluster	16
Task 5.2 – Delete CloudFormation stack.....	17
APPENDIX	19
Appendix 1 – Setting up PuTTY and connecting via SSH	19

PART 1. Setup Environment

Please log into the AWS Management Console and select AWS region as noted below.

You will be using the **Singapore (ap-southeast-1)** region.

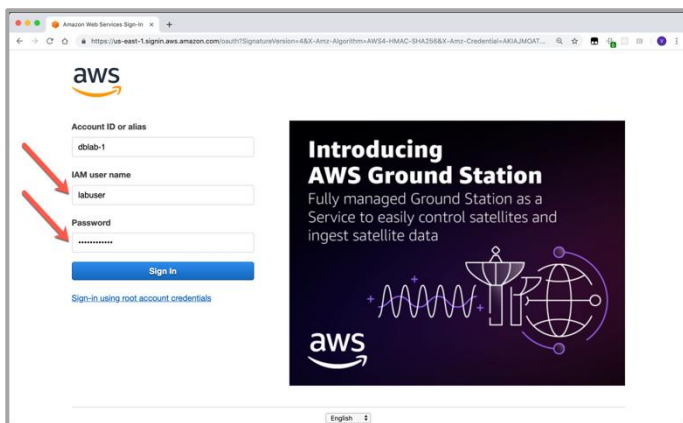
In this part of the lab you will leverage AWS CloudFormation to provision an Aurora MySQL 5.6 compatible database cluster, along with a Linux EC2 instance to be used as a workstation. You will connect to the workstation using SSH.

The environment deployed using CloudFormation includes several components, as listed below. Please download the CloudFormation template (instructions below) and review it for more details.

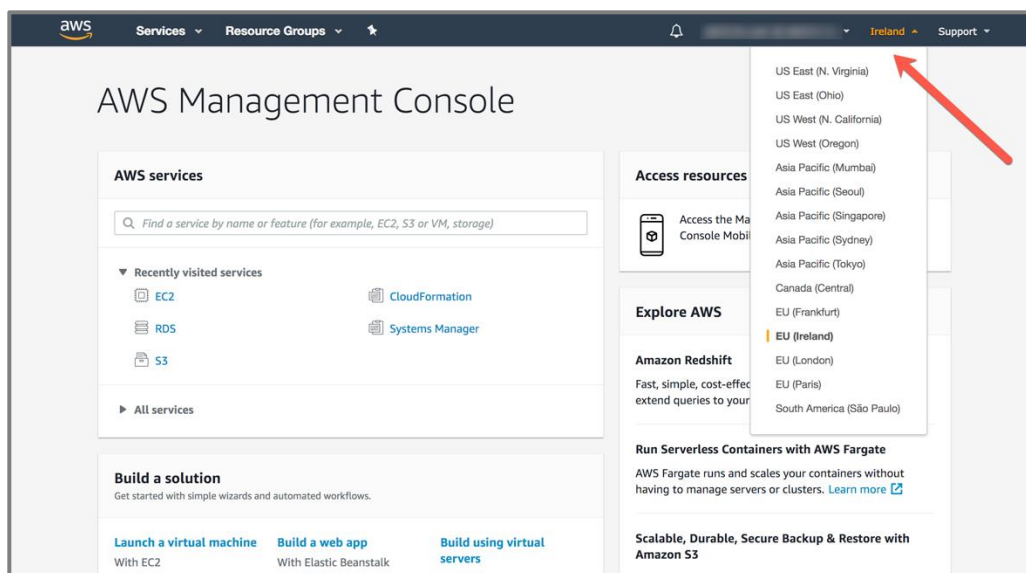
1. Amazon VPC network configuration with public and private subnets
2. Database subnet group and relevant security groups for the cluster and workstation
3. Amazon EC2 instance configured with the software components needed for the lab
4. Roles with access permissions for the workstation and cluster permissions for enhanced monitoring, S3 access and logging
5. Custom cluster and DB instance parameter groups for the Amazon Aurora cluster, enabling logging and performance schema
6. Amazon Aurora DB cluster with 2 nodes: a writer and read replica
7. Read replica auto scaling configuration
8. AWS Systems Manager command document to execute a load test

Task 1.1 – Sign in to Management Console, Select Region and Create Key Pair

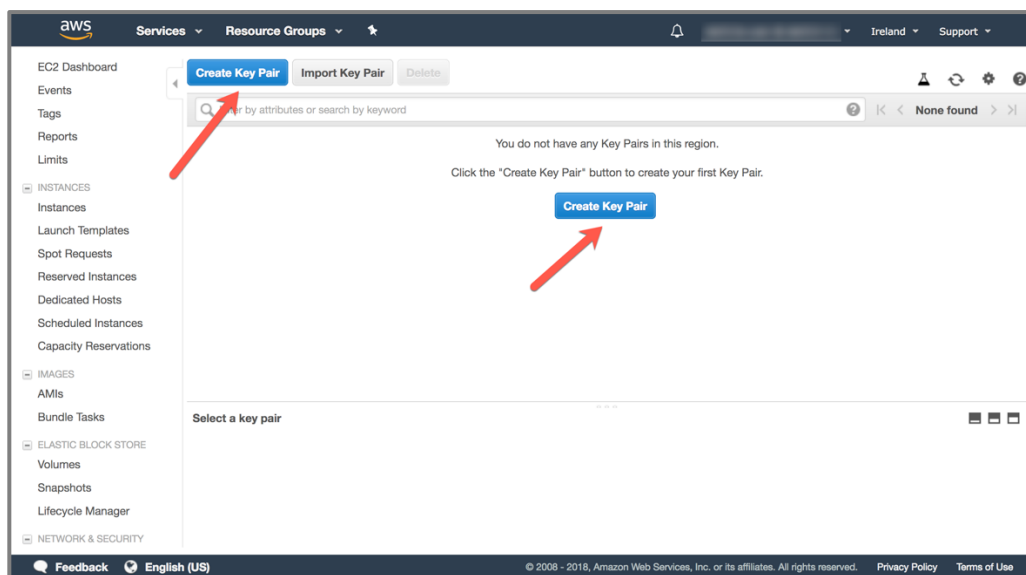
1. **Sign In** to AWS Console.




2. Ensure the **Singapore (ap-southeast-1)** region is selected in the top right corner, if not use that dropdown to choose the correct region



3. Open the **Key Pairs** section of the EC2 service console, using this short link: <https://amzn.to/2IzP387>.
4. Ensure you are still in the correct region, and click **Create Key Pair**.



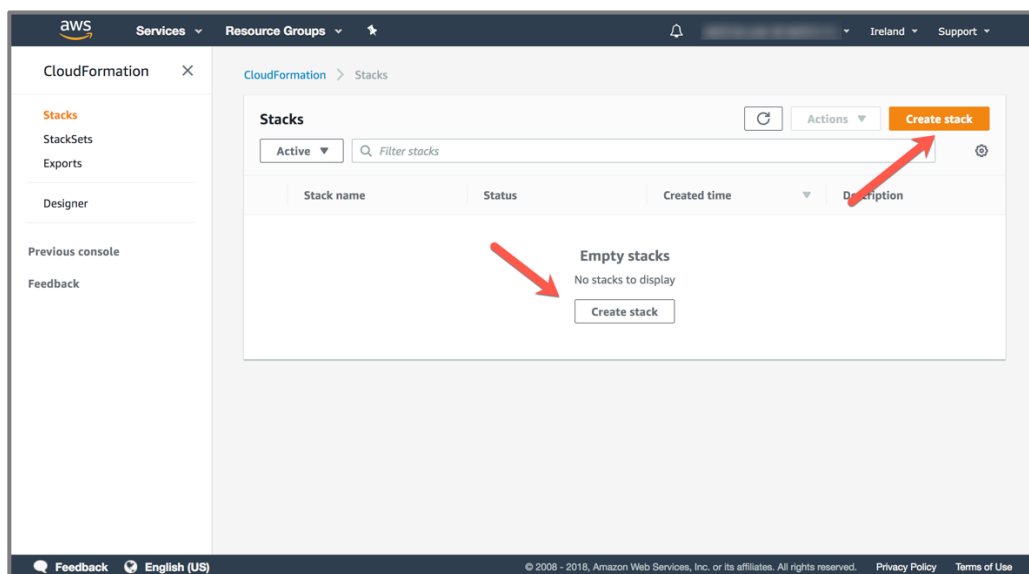
5. Name the key pair "dblabkeys" and then click **Create** and download the file named **dblabkeys.pem** to your computer, save it in a memorable location like your desktop. You will need this file later in the lab.



Task 1.2 - Creating a Stack using CloudFormation

1. Download the CloudFormation template named **lab_template.yml** from <http://bit.ly/aurora-dblab-template>. Save it in a memorable location such as your desktop, you will need to reference it.
2. Open the **CloudFormation** service console located at: <https://amzn.to/2Bln3Jh>.
3. Click **Create Stack**.

Notice: The CloudFormation console has been upgraded recently. Depending on your previous usage of the CloudFormation console UI, you may see the old design or the new design, you may also be presented with a prompt to toggle between them. In this lab we are using the **new design** for reference, although the steps will work similarly in the old console design as well, if you are more familiar with it.



4. Select the radio button named **Upload a template**, then **Choose file** and select the template file you downloaded previously named **lab_template.yml** and then click **Next**.

5. In the field named **Stack Name**, enter the value **“dblabstack”**, select the **ec2KeyPair** value as **“dblabkeys”** (the key pair you have created previously) and then click **Next**.

- On the **Configure stack options** page, leave the defaults as they are, scroll to the bottom and click **Next**.

aws Services Resource Groups Ireland Support

IAM role - optional
Choose the IAM role for CloudFormation to use for all operations performed on the stack.
IAM role na... Sample-role-name Remove

Advanced options

- Stack policy
Defines the resources that you want to protect from unintentional updates during a stack update.
- Rollback configuration
Specify alarms for CloudFormation to monitor when creating and updating the stack. If the operation breaches an alarm threshold, CloudFormation rolls it back. [Learn more.](#)
- Notification options
You can choose an IAM role that CloudFormation uses to create, modify, or delete resources in the stack. If you don't choose a role, CloudFormation uses the permissions defined in your account. [Learn more.](#)
- Stack creation options

Cancel Previous **Next**

Feedback English (US) © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

- On the **Review dblabstack** page, scroll to the bottom, check the box that reads: **I acknowledge that AWS CloudFormation might create IAM resources with custom names** and then click **Create**.

aws Services Resource Groups Ireland Support

Stack creation options

Rollback on failure
Enabled

Timeout
-

Termination protection
Disabled

Quick-create link

Capabilities

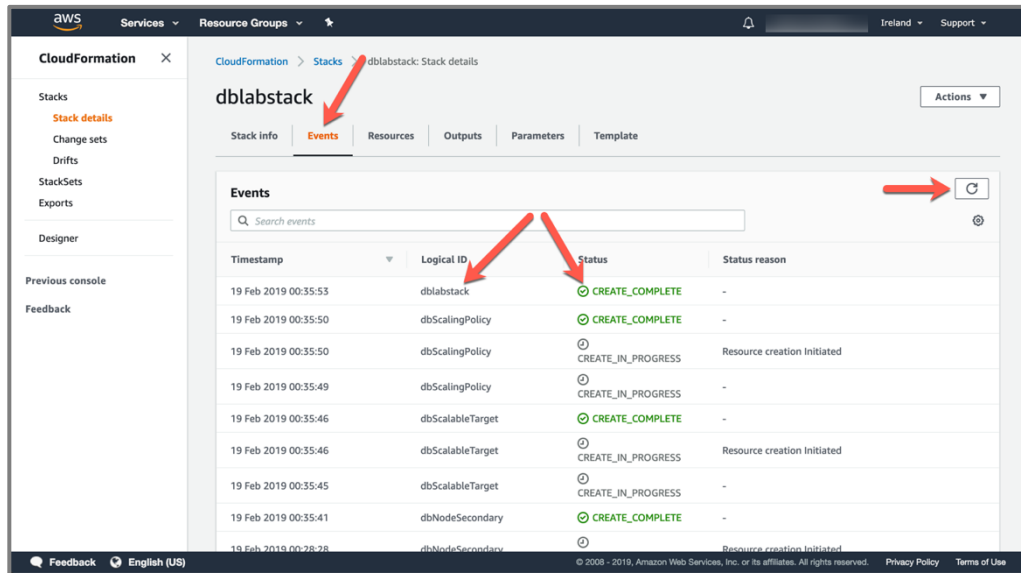
The following resource(s) require capabilities: [AWS::IAM::InstanceProfile, AWS::IAM::Role]
This template contains Identity and Access Management (IAM) resources. Check that you want to create each of these resources and that they have the minimum required permissions. In addition, they have custom names. Check that the custom names are unique within your AWS account. [Learn more.](#)

☒ I acknowledge that AWS CloudFormation might create IAM resources with custom names.

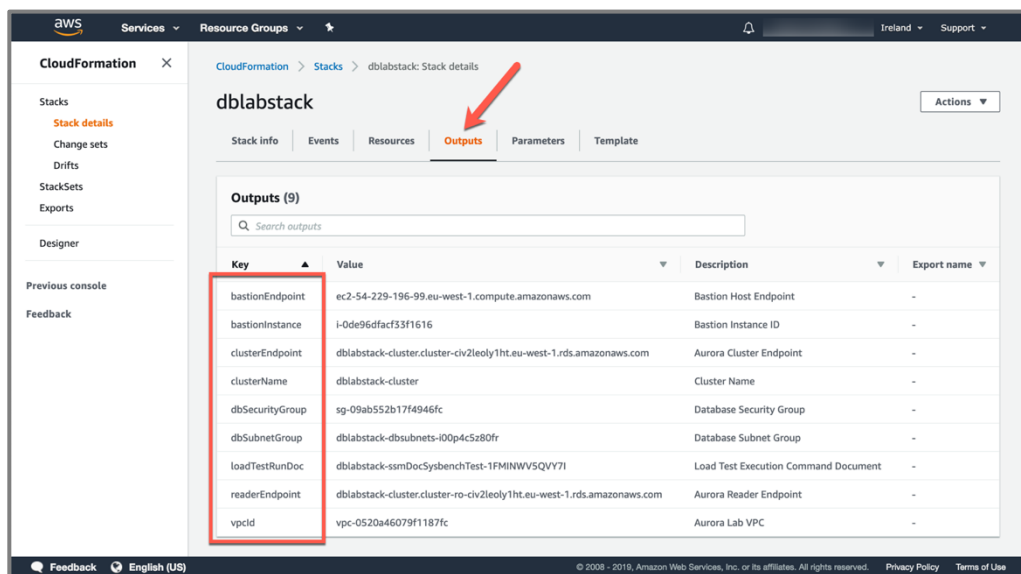
Cancel Previous Create change set **Create stack**

Feedback English (US) © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

- The stack will take approximatively 20 minutes to provision, you can monitor the status on the **Stack detail** page. You can monitor the progress of the stack creation process by refreshing the **Events** tab. The latest event in the list will indicate **CREATE_COMPLETE** for the **dblabstack** resource.



- Once the status of the stack is **CREATE_COMPLETE**, click on the **Outputs** tab. The values here will be critical to the completion of the remainder of the lab. Please take a moment to save these values somewhere that you will have easy access to them during the remainder of the lab. The names that appear in the **Key** column are referenced directly in the instructions in subsequent steps, using the parameter format: **[outputKey]**



Task 1.3 - Connecting to the workstation EC2 instance

For Windows users: We will use PuTTY and PuTTY Key Generator to connect to the workstation using SSH. If you do not have these applications already installed please use the steps in [Appendix 1 - Setting up PuTTY and connecting via SSH](#) below.

For macOS or Linux users: You can connect using the following command from a terminal, however you will need to change the permissions of the certificate file first:

```
chmod 0600 [path to downloaded .pem file]
```

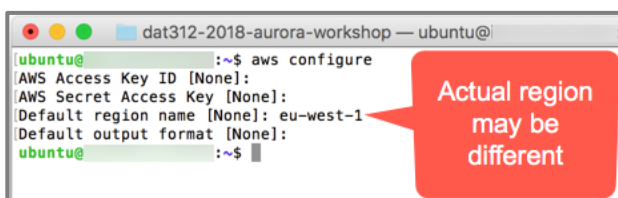
```
ssh -i [path to downloaded .pem file] ubuntu@[bastionEndpoint]
```

Task 1.4 – Set up the AWS CLI and seed the DB Cluster

1. Enter the following command in the SSH console to configure the AWS CLI:

```
aws configure
```

Then select the defaults for everything except the default region name. For the default region name, enter "ap-southeast-1".



2. Connect to the Aurora database using the following command:

```
mysql -h [clusterEndpoint] -u masteruser -p mylab
```

Unless otherwise specified the cluster master username is **masteruser** and the password is **Password1**

3. Run the following queries on the database server, they will create a table, and load data from S3 into it:

```
DROP TABLE IF EXISTS `sbtest1`;

CREATE TABLE `sbtest1` (
  `id` int(10) unsigned NOT NULL AUTO_INCREMENT,
  `k` int(10) unsigned NOT NULL DEFAULT '0',
  `c` char(120) NOT NULL DEFAULT '',
  `pad` char(60) NOT NULL DEFAULT '',
  PRIMARY KEY (`id`),
  KEY `k_1` (`k`)
) ENGINE=InnoDB DEFAULT CHARSET=latin1;
```



```
LOAD DATA FROM S3 MANIFEST
's3-us-west-2://auroraworkshopassets/data/sbtable/sample.manifest'
REPLACE
INTO TABLE sbtest1
CHARACTER SET 'latin1'
FIELDS TERMINATED BY ','
LINES TERMINATED BY '\r\n';
```

PART 2. Cluster Endpoints and Auto Scaling

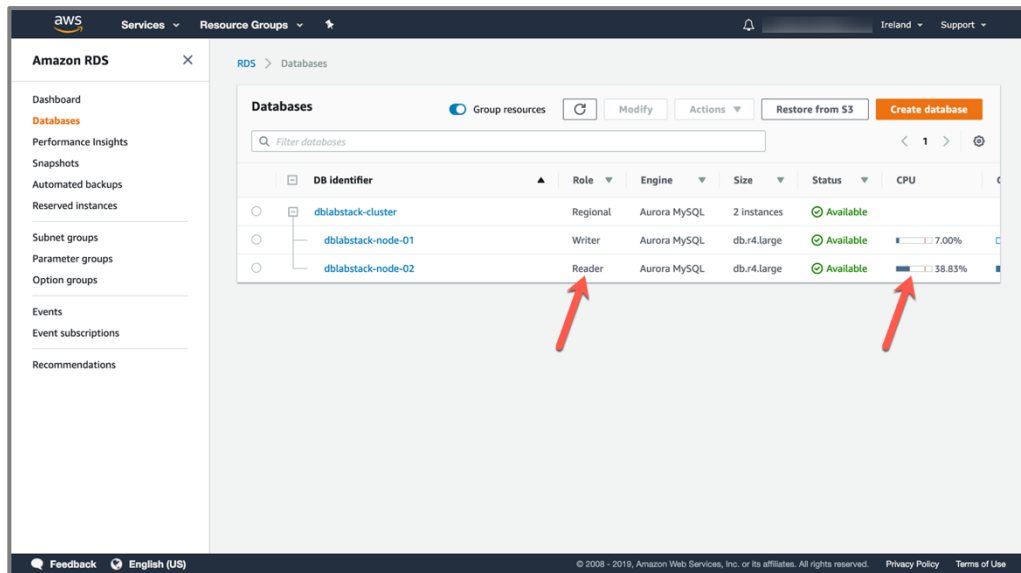
In this part we will explore the cluster endpoints and how auto scaling of read replicas operates.

Task 2.1 – Running a read-only workload

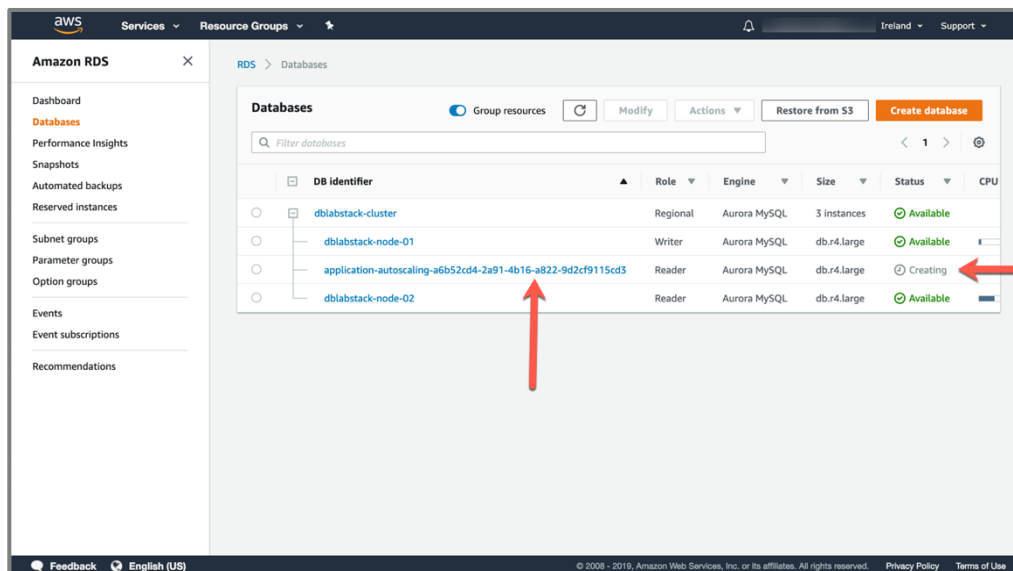
1. On the bastion host, execute the following statement:

```
python loadtest.py [readerEndpoint]
```

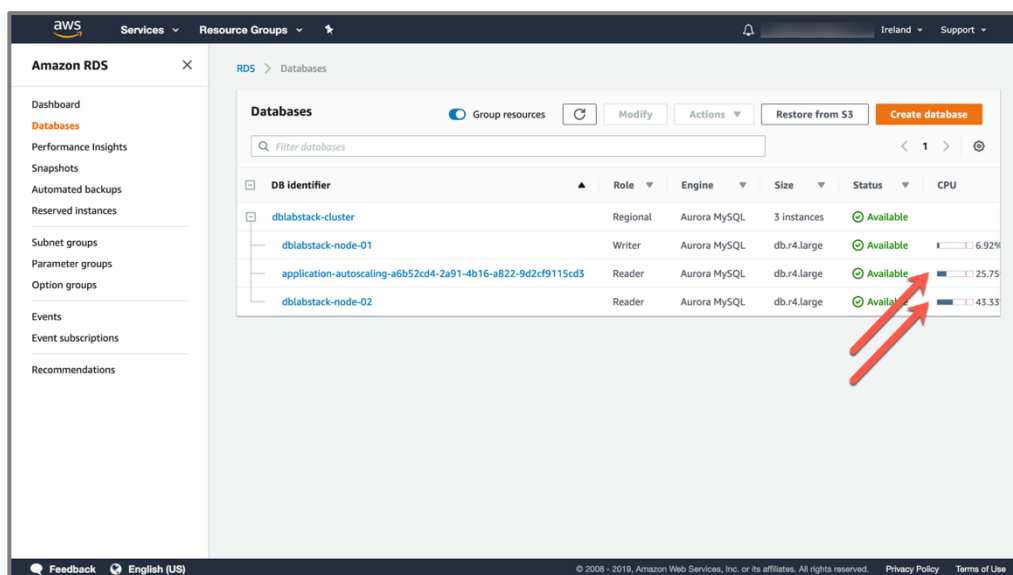
2. Open the **Amazon RDS** service console located at: <https://amzn.to/2GNcMPk>.
3. Take note that the reader node is currently receiving load. It may take a minute or more for the metrics to fully reflect the incoming load.



- After a few minutes return to the list of instances and notice that a new reader is being provisioned to your cluster.



- Once the replicas are added, note that they are starting to receive load.



- You can now type CTRL+C on the bastion host to end the read load, if you wish to. After a while the additional readers will be removed automatically.

PART 3. Cloning and backtracking databases

Task 3.1 - Creating a Clone

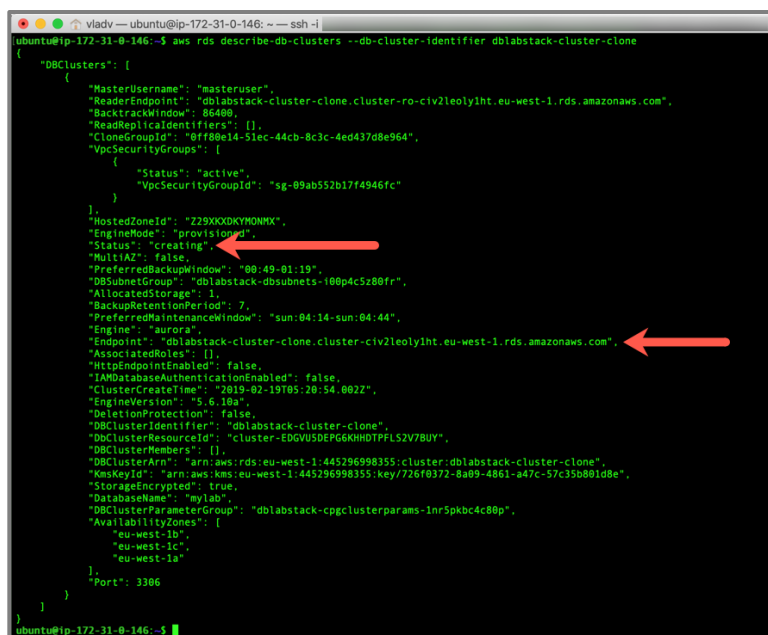
1. On the bastion host, enter:

```
aws rds restore-db-cluster-to-point-in-time --restore-type copy-on-write --
-use-latest-restorable-time --source-db-cluster-identifier [clusterName] --
-db-cluster-identifier [clusterName]-clone --vpc-security-group-ids
[dbSecurityGroup] --db-subnet-group-name [dbSubnetGroup] --backtrack-
window 86400
```

2. Next, to check the status of the creation of your clone, enter the following command on the bastion host. The cloning process can take several minutes to complete. See the example output below.

Note: This step will create the Aurora DB cluster itself, but without any compute nodes. You will add a computer node in the next step.

```
aws rds describe-db-clusters --db-cluster-identifier [clusterName]-clone
```



```
ubuntu@ip-172-31-0-146:~$ aws rds describe-db-clusters --db-cluster-identifier dblabstack-cluster-clone
{
  "DBClusters": [
    {
      "MasterUsername": "masteruser",
      "ReaderEndpoint": "dblabstack-cluster-clone.cluster-ro-civ2leolyint.eu-west-1.rds.amazonaws.com",
      "BacktrackWindow": 86400,
      "ReadReplicaIdentifiers": [],
      "CloneGroupId": "9ff8de14-51ec-44cb-8c3c-4ed437d8e964",
      "VpcSecurityGroups": [
        {
          "Status": "active",
          "VpcSecurityGroupId": "sg-09ab552b17f4946fc"
        }
      ],
      "HostedZoneId": "Z29KXQXKYNOMMX",
      "EngineMode": "provisioner",
      "Status": "creating",
      "MultiAZ": false,
      "PreferredBackupWindow": "00:49-01:19",
      "DBSubnetGroup": "dblabstack-dbsubnets-100p4c5z80fr",
      "AllocatedStorage": 1,
      "BackupRetentionPeriod": 7,
      "PreferredMaintenanceWindow": "sun:04:14-sun:04:44",
      "Engine": "aurora",
      "Endpoint": "dblabstack-cluster-clone.cluster-ro-civ2leolyint.eu-west-1.rds.amazonaws.com",
      "AssociatedRoles": [],
      "HttpEndpointEnabled": false,
      "IAMDatabaseAuthenticationEnabled": false,
      "ClusterCreateTime": "2019-02-19T05:20:54.002Z",
      "EngineVersion": "5.6.10a",
      "DeletionProtection": false,
      "DBClusterIdentifier": "dblabstack-cluster-clone",
      "DBClusterResourceId": "cluster-EDGVUSDEPG6KHHDTPFLS2V78UY",
      "DBClusterMembers": [],
      "DBClusterArn": "arn:aws:rds:eu-west-1:445296998355:cluster:dblabstack-cluster-clone",
      "KeyId": "arn:aws:kms:eu-west-1:445296998355:key/726f0372-8a09-4861-a47c-57c35b0010be",
      "StorageEncrypted": true,
      "DatabaseName": "mylab",
      "DBClusterParameterGroup": "dblabstack-cpgclusterparams-1nr5pkbc4c80p",
      "AvailabilityZones": [
        "eu-west-1b",
        "eu-west-1c",
        "eu-west-1a"
      ],
      "Port": 3306
    }
  ]
}
```

3. Take note of both the **"Status"** and the **"Endpoint."** Once the **Status** becomes **available**, you can add an instance to the cluster and once the instance is added, you will want to connect to the cluster via the **Endpoint** value. To add an instance to the cluster once the status becomes **available**, enter the following:

```
aws rds create-db-instance --db-instance-class db.r4.large --engine aurora
--db-cluster-identifier [clusterName]-clone --db-instance-identifier
[clusterName]-clone-instance
```

4. To check the creation of the instance, enter the following at the command line:

```
aws rds describe-db-instances --db-instance-identifier [clusterName]-clone-instance
```

- Once the **DBInstanceStatus** changes from **creating** to **available**, you have a functioning clone. Creating a node in a cluster also takes several minutes.

```

{
  "DBInstances": [
    {
      "PubliclyAccessible": false,
      "MasterUsername": "masteruser",
      "MonitoringInterval": 0,
      "LicenseModel": "general-public-license",
      "VpcSecurityGroups": [
        {
          "Status": "active",
          "VpcSecurityGroupId": "sg-09ab552b17f4946fc"
        }
      ],
      "CopyTagsToSnapshot": false,
      "OptionGroupMemberships": [
        {
          "Status": "in-sync",
          "OptionGroupName": "default:aurora-5.6"
        }
      ],
      "PendingModifiedValues": {},
      "Engine": "aurora",
      "MultiAZ": false,
      "DBSecurityGroups": [],
      "DBParameterGroups": [
        {
          "DBParameterGroupName": "default:aurora5.6",
          "ParameterApplyStatus": "in-sync"
        }
      ],
      "PerformanceInsightsEnabled": false,
      "AutoMinorVersionUpgrade": true,
      "PreferredBackupWindow": "00:49-01:19",
      "PromotionTier": 1,
      "DBSubnetGroup": {
        "Subnets": [
          {
            "SubnetStatus": "Active",
            "SubnetIdentifier": "subnet-0652425f9a24f33aa",
            "SubnetAvailabilityZone": {
              "Name": "eu-west-1c"
            }
          },
          {
            "SubnetStatus": "Active",
            "SubnetIdentifier": "subnet-0f31f662b9ba7fc8",
            "SubnetAvailabilityZone": {
              "Name": "eu-west-1b"
            }
          },
          {
            "SubnetStatus": "Active",
            "SubnetIdentifier": "subnet-0e1e540a0ea85bde",
            "SubnetAvailabilityZone": {
              "Name": "eu-west-1a"
            }
          }
        ],
        "DBSubnetGroupName": "dblabstack-dbsubnets-100p4c5z80fr",
        "VpcId": "vpc-0520a46079f1187fc",
        "DBSubnetGroupDescription": "dblabstack-db-subnet-group",
        "SubnetGroupStatus": "Complete"
      },
      "ReadReplicaDBInstanceIdentifiers": [],
      "AllocatedStorage": 1,
      "DBInstanceArn": "arn:aws:rds:eu-west-1:445296998355:db:dblabstack-cluster-clone-instance",
      "BackupRetentionPeriod": 7,
      "DBName": "mylab",
      "PreferredMaintenanceWindow": "mon:02:06-mon:02:36",
      "DBInstanceStatus": "creating",
      "IAMDatabaseAuthenticationEnabled": false,
      "EngineVersion": "5.6.10a",
      "DeletionProtection": false,
      "AvailabilityZone": "eu-west-1c",
      "DomainMemberships": [],
      "DBClusterIdentifier": "dblabstack-cluster-clone",
      "StorageType": "aurora",
      "DbResourceId": "db-v43TLNU5FLUKQWZYSYVHFN3JXI",
      "CACertificateIdentifier": "rds-ca-2015",
      "KmsKeyId": "arn:aws:kms:eu-west-1:445296998355:key/726f0372-8a09-4861-a47c-57c35b801d8e",
      "StorageEncrypted": true,
      "AssociatedRoles": [],
      "DBInstanceClass": "db.r4.large",
      "DBInstancePort": 0,
      "DBInstanceIdentifier": "dblabstack-cluster-clone-instance"
    }
  ]
}

```

- Once your instance is created, connect to the instance using the following command:

```
mysql -h [cluster endpoint of clone cluster] -u masteruser -p mylab
```

Note: the master user account credentials will be the same as with the source of the cloned cluster. If you customized the CloudFormation template and changed the values, use the customized username and password.

- In order to verify that the clone is identical to the source, we will perform a checksum of the sbtest1 table using the following:

```
checksum table sbtest1;
```

8. The output of your commands should look similar to the example below:

```
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 5.6.10 MySQL Community Server (GPL)


Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MySQL [(none)]> use sysbench;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
MySQL [sysbench]> checksum table sbtest1;
+-----+-----+
| Table          | Checksum |
+-----+-----+
| sysbench.sbtest1 | 2321424769 |
+-----+-----+
1 row in set (15.25 sec)

MySQL [sysbench]> 
```



9. Please take note of the value for your specific clone cluster.
10. Next, we will disconnect from the clone and connect to the original cluster with the following:

quit;

mysql -h [clusterEndpoint] -u masteruser -p mylab

11. Next, we will execute the same commands that we executed on the clone:

checksum table sbtest1;

12. Please take note of the value for your specific source cluster. The checksum should be identical.

```
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 5.6.10 MySQL Community Server (GPL)


Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MySQL [(none)]> use sysbench;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
MySQL [sysbench]> checksum table sbtest1;
+-----+-----+
| Table          | Checksum |
+-----+-----+
| sysbench.sbtest1 | 2321424769 |
+-----+-----+
1 row in set (15.25 sec)

MySQL [sysbench]> 
```



Task 3.2 - Backtracking the Database

1. Reconnect to the cloned cluster using:

```
quit;

mysql -h [cluster endpoint of clone cluster] -u masteruser -p mylab
```

2. Drop the **sbtest1** table:

Note: Consider executing the commands below one at a time, waiting a few seconds between each one. This will make it easier to determine a good point in time for testing backtrack.

```
select current_timestamp();

drop table sbtest1;

select current_timestamp();

quit;
```

3. Remember or save the time markers displayed above, you will use them as references later.
4. Run the following command to replace the dropped table using the sysbench command:

```
sysbench oltp_write_only --threads=1 --mysql-host=[cluster endpoint of clone cluster] --mysql-user=masteruser --mysql-password=Password1 --mysql-port=3306 --tables=1 --mysql-db=mylab --table-size=1000000 prepare
```

5. Reconnect to the cloned cluster, and checksum the table again, the checksum value should be different than both the original clone value and source cluster:

```
mysql -h [cluster endpoint of clone cluster] -u masteruser -p mylab

checksum table sbtest1;

quit;
```

6. Backtrack the database to a time slightly after the second time marker. (Right after dropping the table).

For e.g. If the second time marker is "2019-06-03 19:19:40", specify *backtrack-to* parameter value as "2019-06-03T19:19:40Z" in the command below.

```
aws rds backtrack-db-cluster --db-cluster-identifier [clusterName] --clone --backtrack-to "yyyy-mm-ddThh:mm:ssZ"
```

7. Run the below command to track the progress of the backtracking operation. The operation should complete in a few minutes. Status will change from "backtracking" to "available" when backtrack is complete.

```
aws rds describe-db-clusters --db-cluster-identifier [clusterName] --clone | grep -i EngineMode -A 2 | grep Status
```

```
ubuntu@ip-172-31-0-100:~$ aws rds describe-db-clusters --db-cluster-identifier dlabstack-cluster | grep -i EngineMode -A 2 | grep Status
  "Status": "backtracking",
ubuntu@ip-172-31-0-100:~$ aws rds describe-db-clusters --db-cluster-identifier dlabstack-cluster | grep -i EngineMode -A 2 | grep Status
  "Status": "available",
```

- Wait for backtracking to complete and then connect back to the database. The **sbtest1** table should be missing from the database.

```
mysql -h [cluster endpoint of clone cluster] -u masteruser -p mylab
show tables;
quit;
```

- Now backtrack again to a time slightly before the first time marker above. (Right before dropping the table).

```
aws rds backtrack-db-cluster --db-cluster-identifier [clusterName] --clone -
-backtrack-to "yyyy-mm-ddThh:mm:ssZ"
```

- Run the below command to track the progress of the backtracking operation. The operation should complete in a few minutes.

```
aws rds describe-db-clusters --db-cluster-identifier [clusterName] --clone |
grep -i EngineMode -A 2 | grep Status
```

- Connect back to the database. The **sbtest1** table should now be available in the database again, but contain the original data set.

```
mysql -h [cluster endpoint of clone cluster] -u masteruser -p mylab
show tables;
quit;
```

PART 4. Performance Insights

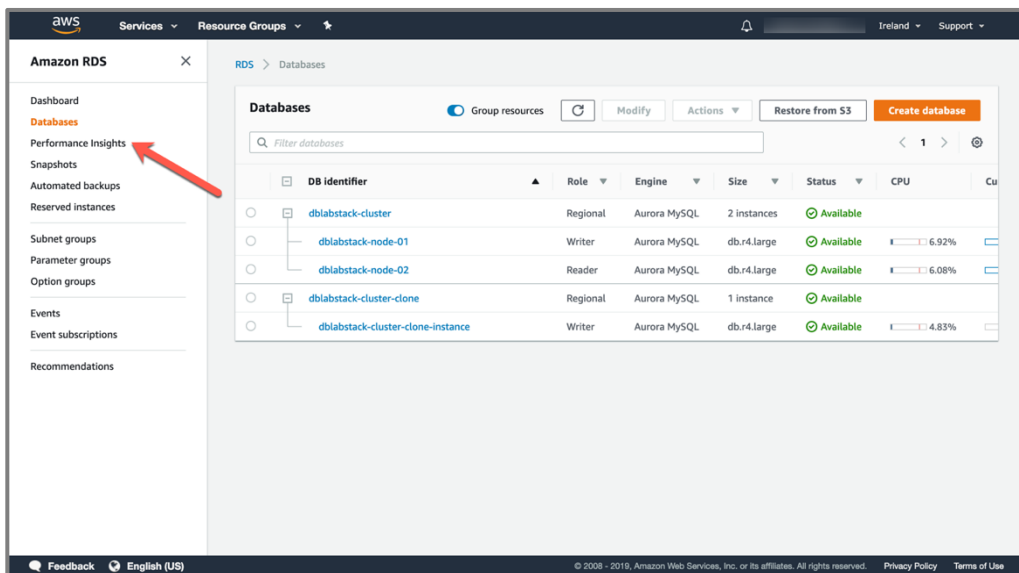
Task 4.1 – Generate Load on Your Database Cluster

- You will use [Percona's TPCC-like benchmark script](#) based on sysbench to generate load. For simplicity we have packaged the correct set of commands in an AWS Systems Manager Command Document. You will use AWS Systems Manager Run Command to execute the test.
- On the workstation host, execute the following statement:

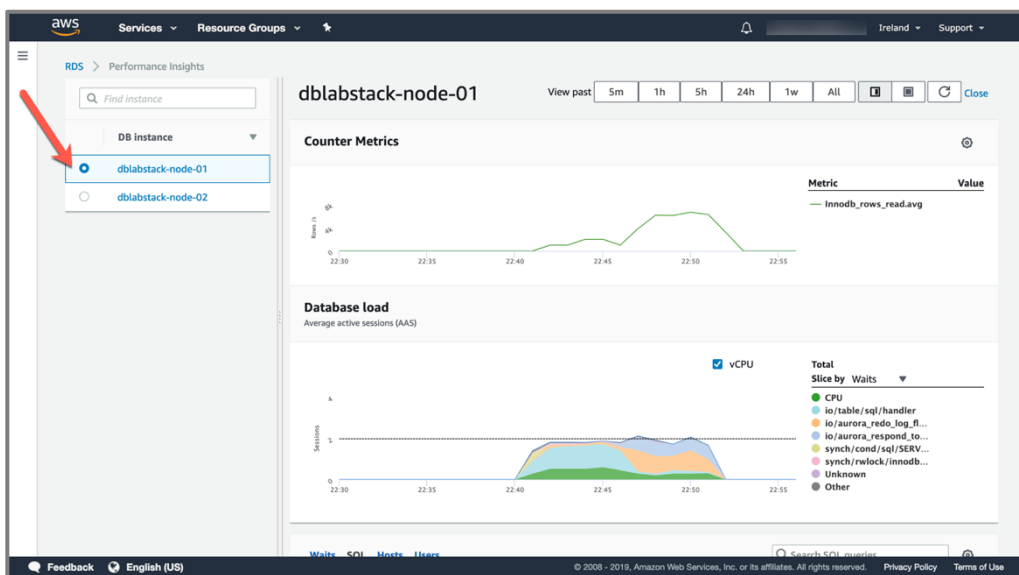
```
aws ssm send-command --document-name [loadTestRunDoc] --instance-ids
[bastionInstance]
```

- The command will be sent to the workstation EC2 instance which will prepare the test data set and run the load test. It may take up to a minute for CloudWatch to reflect the additional load in the metrics.

- Navigate to the RDS service console (<https://amzn.to/2GNcMPk>) and click on **Performance Insights** in the left side navigation bar.



- Examine the performance of your DB instance **demostack-node-01** using Performance Insights. What conclusions can you reach?



Part 5. Delete AWS resources

Follow these steps to clean up AWS resources after you are done with all the lab activities.

Task 5.1 – Delete Cloned Aurora Cluster

- Delete the Aurora instance on the cloned cluster by entering the following on the bastion host.

```
aws rds delete-db-instance --db-instance-identifier [clusterName]-clone-instance --skip-final-snapshot
```


2. Verify the Aurora instance is deleted, by entering the following on the command line.

```
aws rds describe-db-instances --db-instance-identifier [clusterName]-clone-instance | grep "DBInstanceStatus"
```

Once the instance is deleted, you will see the following error message while running the command.

```
ubuntu@ip-172-31-0-188:~$ aws rds describe-db-instances --db-instance-identifier dlabstack-cluster-clone-instance | grep "DBInstanceStatus"
An error occurred (DBInstanceNotFound) when calling the DescribeDBInstances operation: DBInstance dlabstack-cluster-clone-instance not found.
ubuntu@ip-172-31-0-188:~$
```

3. Delete the cloned Aurora cluster by entering the following on the command line.

```
aws rds delete-db-cluster --db-cluster-identifier [clusterName]-clone -skip-final-snapshot
```

4. Verify the Cloned Aurora Cluster is delete, by entering the following on the command line.

```
aws rds describe-db-clusters --db-cluster-identifier [clusterName]-clone | grep "Status"
```

When the cluster is deleted, you will see the following error message while running the command.

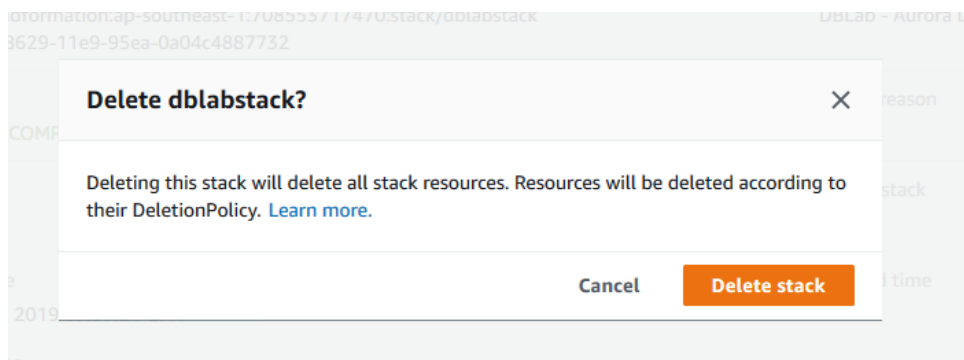
```
ubuntu@ip-172-31-0-188:~$ aws rds describe-db-clusters --db-cluster-identifier dlabstack-cluster-clone | grep "Status"
An error occurred (DBClusterNotFoundFault) when calling the DescribeDBClusters operation: DBCluster dlabstack-cluster-clone not found.
ubuntu@ip-172-31-0-188:~$
```

Task 5.2 – Delete CloudFormation stack

1. Open the **CloudFormation** service console located at: <https://amzn.to/2Bln3Jh>.
2. Select the CloudFormation stack you created as part of this lab in the left panel and select Delete button on the right side of the screen.

The screenshot shows the AWS CloudFormation console. In the left-hand navigation pane, under 'Stacks (1)', the 'dblabstack' stack is listed with a status of 'CREATE_COMPLETE'. A red arrow points to this stack entry. In the main panel, the 'dblabstack' stack is selected, and the 'Stack info' tab is active. In the top right corner of this panel, there are buttons for 'Delete', 'Update', and 'Stack actions'. A red arrow points to the 'Delete' button.

3. Select Delete stack when prompted. It will take 15 minutes to delete the stack.



APPENDIX

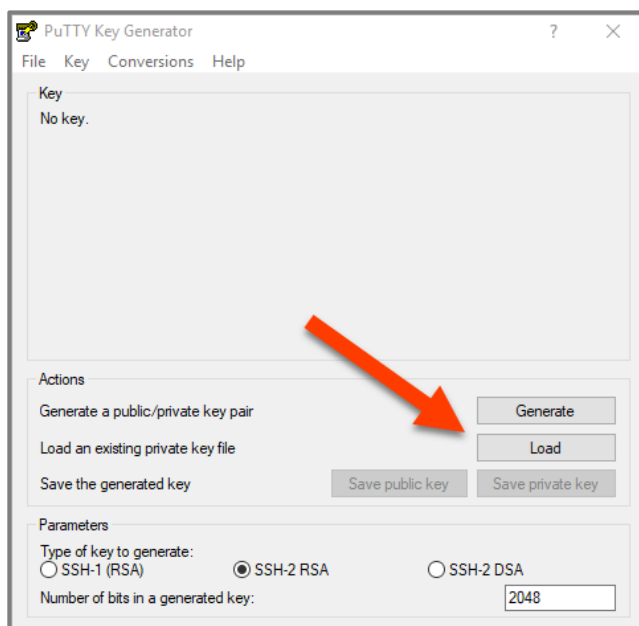
Appendix 1 – Setting up PuTTY and connecting via SSH

For Windows users, please download **PuTTY** (putty) and the **PuTTY Key Generator** (puttygen) from the following links:


<https://bit.ly/2ETxptZ>

<https://bit.ly/2AEoOHp>

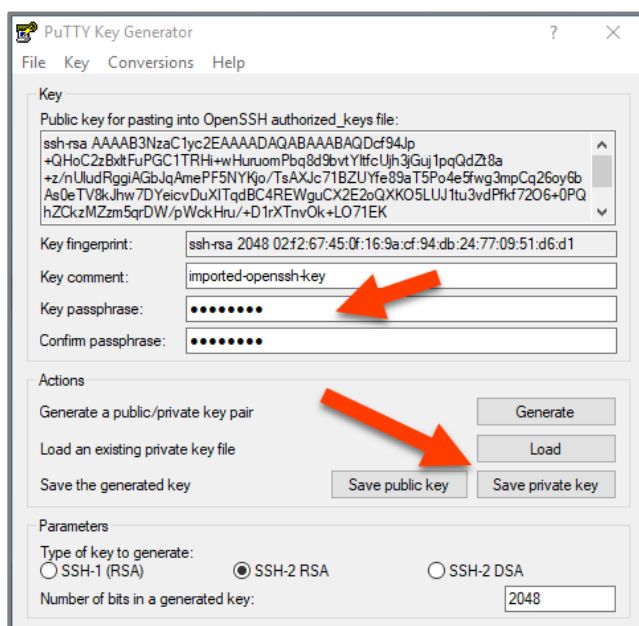
1. Once you have downloaded putty and puttygen, open **puttygen** and click on **Load**.



2. Please make sure that the file filter is set to "All Files (*.*)" and then select **dblackkeys.pem**.

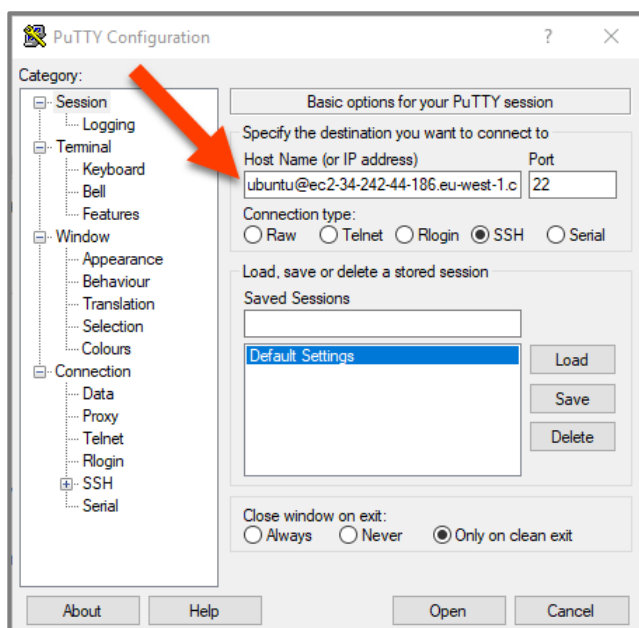


- Fill in the **Key passphrase** and **Confirm passphrase** fields with a password of your choice that will be used to encrypt your private key and then click **Save private key**. Please use "db1abkeys.ppk" as your new key name.

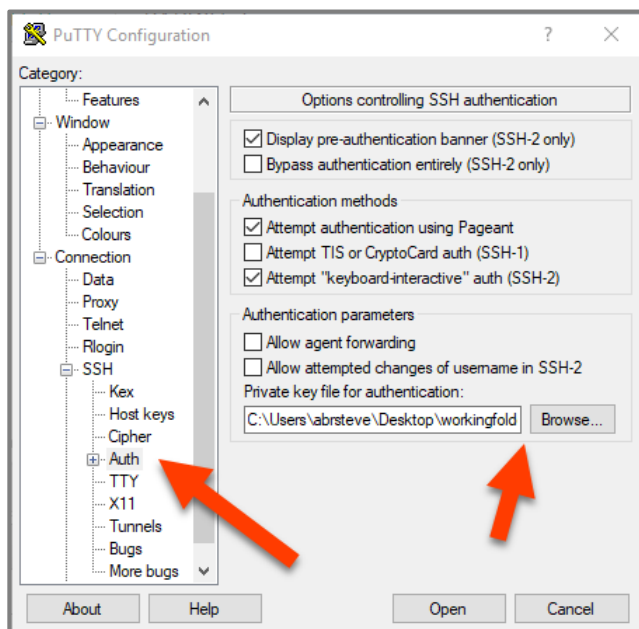


- Next, open putty and enter into the **Host Name (or IP address)** field the following value:

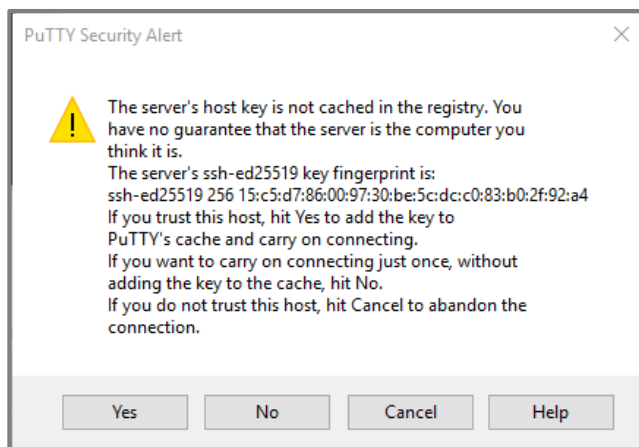
ubuntu@[bastionEndpoint]



- Next, navigate within PuTTY to **Connection** → **SSH** → **Auth** and browse to the **reinventdemo.ppk** file that you created with the PuTTY Key Generator previously, and then click **Open**.



- When prompted by the PuTTY Security Alert, click **Yes**.



- Next, enter the password that you configured when you created the **dblabkeys.ppk** private file previously.

```
Using username "ubuntu".
Authenticating with public key "imported-openssh-key"
Passphrase for key "imported-openssh-key":
```